

# REAL-TIME DIGITAL SYSTEMS DESIGN AND VERIFICATION WITH FPGAS

## ECE 387 – LECTURE 5

PROF. DAVID ZARETSKY

DAVID.ZARETSKY@NORTHWESTERN.EDU

# AGENDA

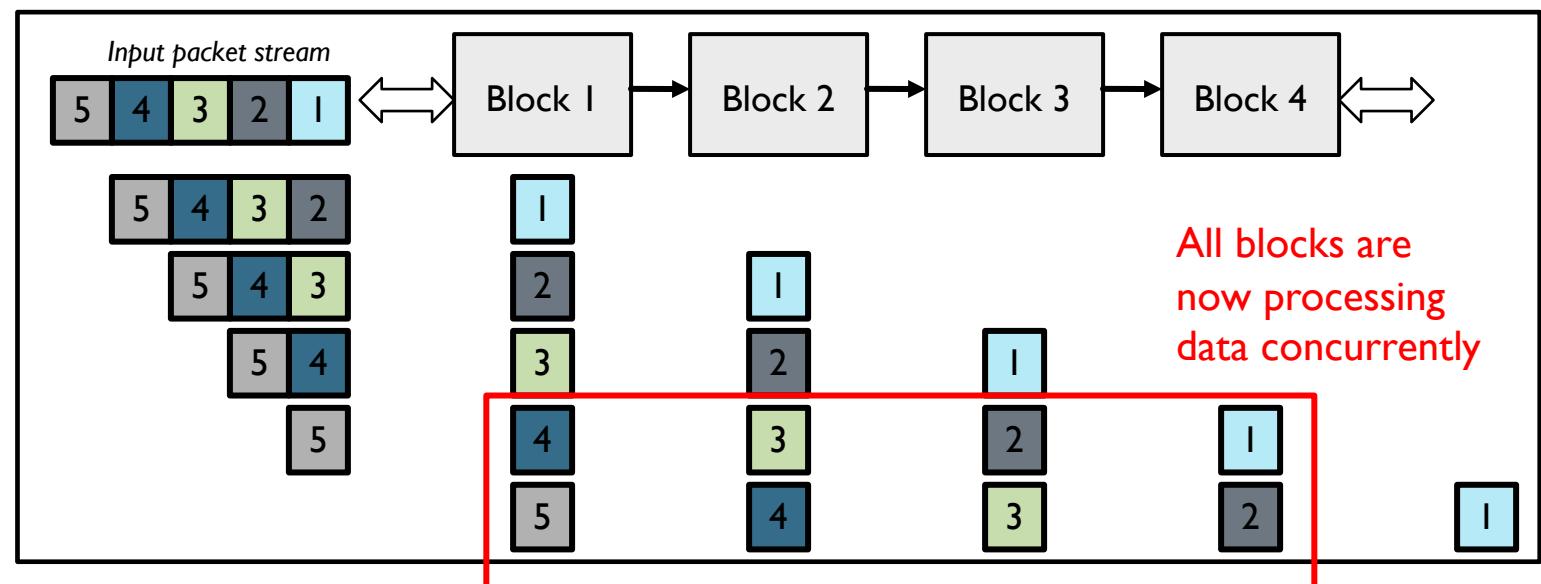
- Streaming Architectures
- FIFOs

# STREAMING ARCHITECTURES IN FPGAS

- FPGAs are unique among programmable platforms in that streaming architectures can be implemented with elements of varying granularity.
- This enables different actors in a dataflow specification to be easily implemented at matched data rates, despite differing computational loads.

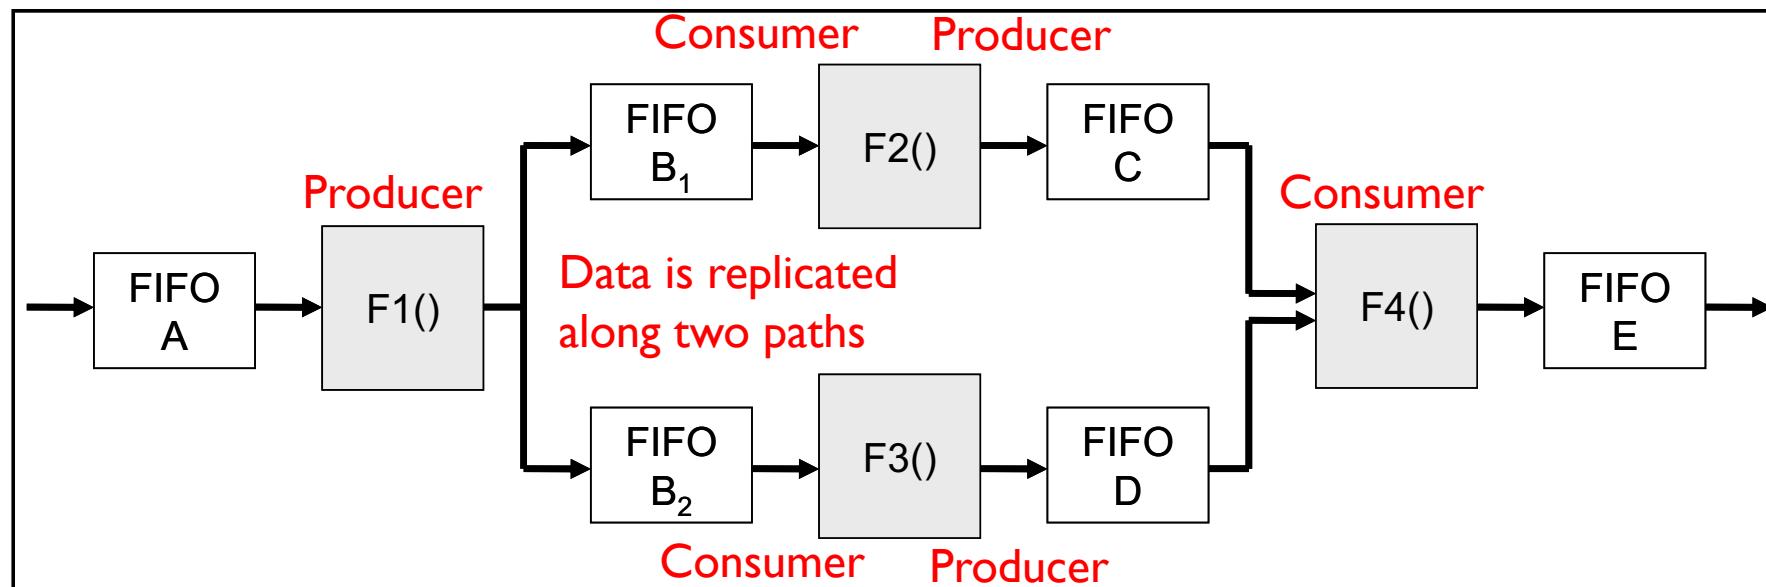
# BLOCK-LEVEL PIPELINING

- Blocks or modules operate on data packets in chunks.
- Blocks are pipelined to allow faster processing
- There is an initial cost in propagating data through the pipeline until all blocks are working concurrently.

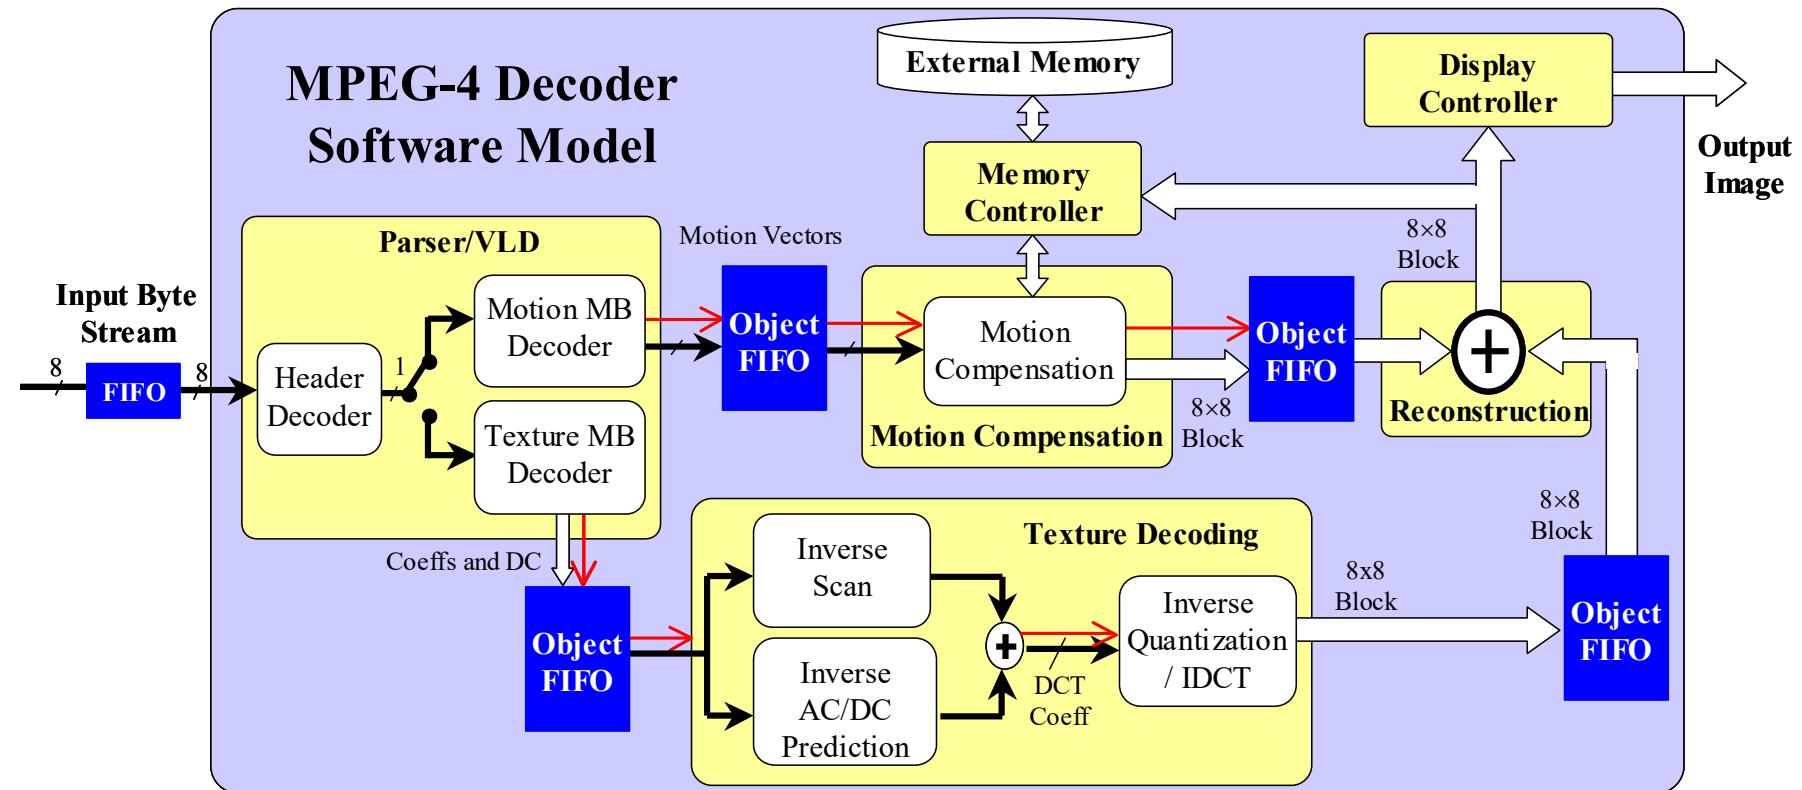


# STREAMING ARCHITECTURES

- Using FIFOs, data is passed between function blocks in bursts, thus allowing blocks to be pipelined and processed concurrently.
- FIFOs can be used when the producer and consumer functions have the same data read and write schemes.

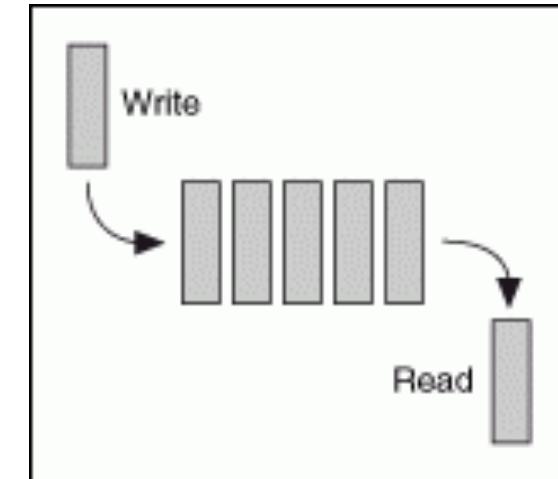


# STREAMING MPEG-4 EXAMPLE



# FIFOS

- A FIFO is a data structure that holds elements in the order they are received
- It provides access to those elements using a first-in, first-out access policy.
- Use FIFOs to transfer data in the following ways:
  - Between parallel processes within one clock domain
  - Across clock domains
  - Arrays passed between components (entities)
  - Between the host computer and the FPGA
  - Between peripherals and the FPGA



# FIFO VS. BLOCK RAMS

- Block RAMs
  - Access data in any order
  - Requires one cycle latency to read cycles
  - Write on demand
  - Storage for full or chunk of data
  - Requires memory addressing
  - Unknown if RAM is empty/full
  - Multi-clock domains can be tricky
  - Useful for complex read/write patterns
- FIFOs
  - Access data sequentially only
  - Data read/write on demand
  - Storage for a small subset of data
  - Does not require complex addressing operations
  - Synchronization signals tell if FIFO is empty/full
  - Can pass data between multi-clock domains
  - Useful for sequential read/write patterns or if producer and consumer patterns are identical

# FIFO MODEL

- FIFOs provide the glue-logic for streaming architectures
- Allow data to pass between different time domains
- Single-cycle read/write – data is instantaneously available
- Eliminates read and write addresses that consume a lot of resources
- If not built correctly, timing issues may emerge

```
module fifo #(
    parameter FIFO_DATA_WIDTH = 32,
    parameter FIFO_BUFFER_SIZE = 1024
)
(
    input  logic reset,
    input  logic wr_clk,
    input  logic wr_en,
    input  logic [FIFO_DATA_WIDTH-1:0] din,
    output logic full,
    input  logic rd_clk,
    input  logic rd_en,
    output logic [FIFO_DATA_WIDTH-1:0] dout,
    output logic empty
);
```

# FIFO READ & WRITE OPERATIONS

- FIFO Write Operations

- If full == 0, data may be written to din
- Data may be written instantaneously (same cycle)
- Write enable writes the data and advances the address pointer for the next write operation

- FIFO Read Operations

- If empty == 0, dout is valid for reading
- Data may be acquired instantaneously (same cycle)
- Read enable advances the address pointer for the next read operation



# FIFO WRITE

```
always_ff @(posedge wr_clk)
begin : p_write_buffer
    if ( (wr_en == 1'b1) && (full_t == 1'b0) ) begin
        fifo_buf[$unsigned(wr_addr[FIFO_ADDR_WIDTH-2:0])] <= din;
    end
end

always_ff @(posedge wr_clk, posedge reset)
begin : p_wr_addr
    if ( reset == 1'b1 )
        wr_addr <= '0;
    else
        wr_addr <= wr_addr_t;
end

assign wr_addr_t = (wr_en == 1'b1 && full_t == 1'b0) ? ($unsigned(wr_addr) + 'h1) : wr_addr;

assign full_t = (wr_addr[FIFO_ADDR_WIDTH-2:0] == rd_addr[FIFO_ADDR_WIDTH-2:0]) &&
                (wr_addr[FIFO_ADDR_WIDTH-1] != rd_addr[FIFO_ADDR_WIDTH-1]) ? 1'b1 : 1'b0;

assign full = full_t;
```

# FIFO READ

```
always_ff @(posedge rd_clk)
begin : p_rd_buffer
    dout <= fifo_buf[$unsigned(rd_addr_t[FIFO_ADDR_WIDTH-2:0])];
end

always_ff @(posedge rd_clk, posedge reset)
begin : p_rd_addr
    if (reset == 1'b1)
        rd_addr <= '0;
    else
        rd_addr <= rd_addr_t;
end

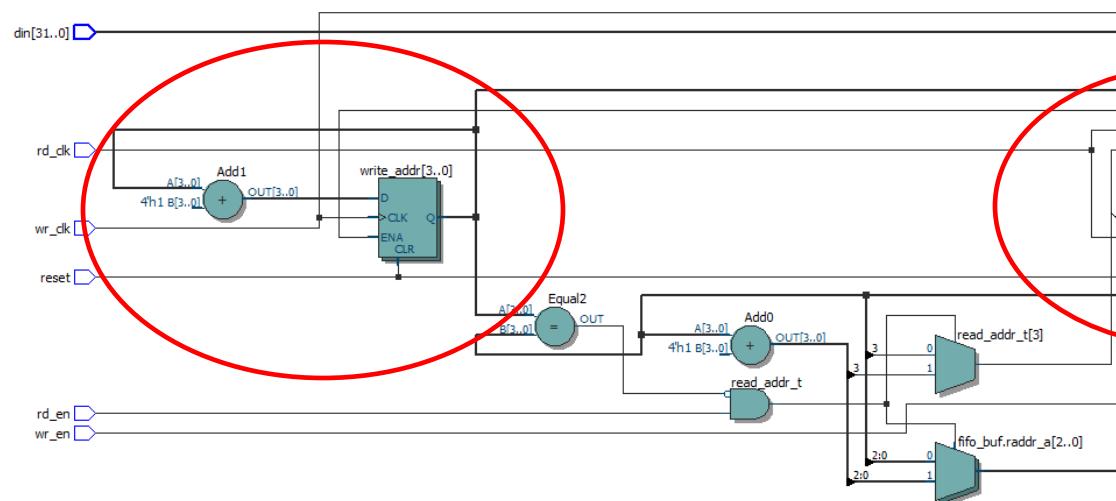
always_ff @(posedge rd_clk, posedge reset)
begin : p_empty
    if (reset == 1'b1)
        empty <= '0;
    else
        empty <= (wr_addr == rd_addr_t) ? 1'b1 : 1'b0;
end

assign rd_addr_t = (rd_en == 1'b1 && empty_t == 1'b0) ? ($unsigned(rd_addr) + 'h1) : rd_addr;
assign empty_t = (wr_addr == rd_addr) ? 1'b1 : 1'b0;
```

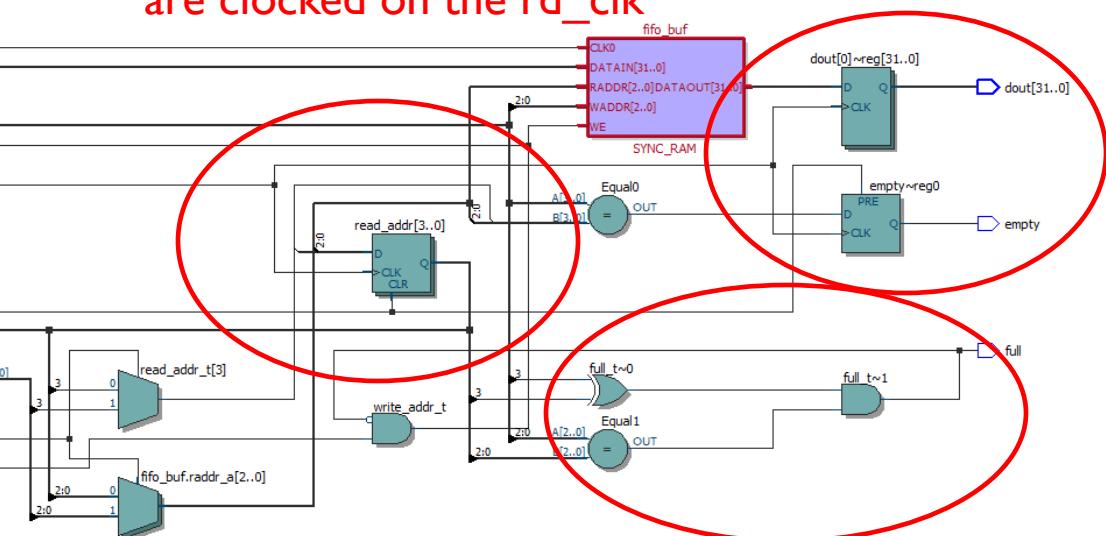
# FIFO TECHNOLOGY VIEW

- Read and Write signals are each registered on different clock domains

Write address is clocked  
on the wr\_clk



Read address, data and empty  
are clocked on the rd\_clk



Full is based on clocked read  
and write addresses

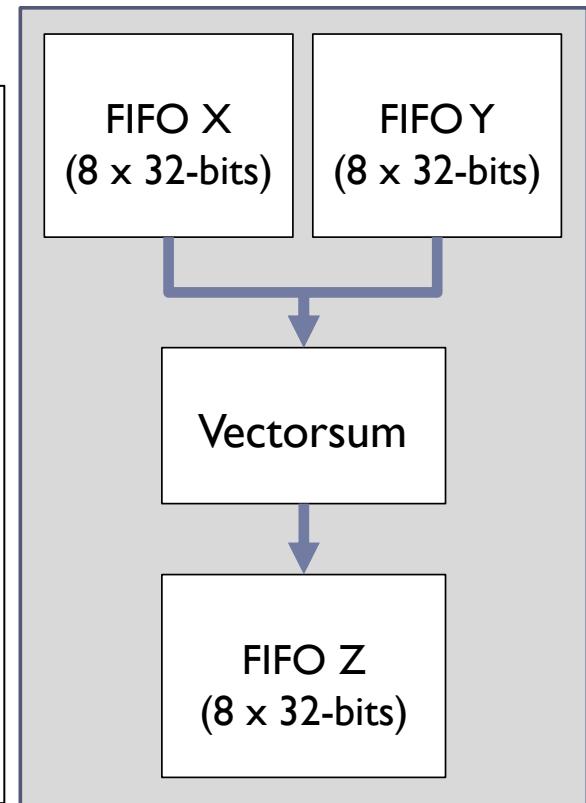
# STREAMING VECTORSUM

- Read arrays X and Y from FIFOs
- Compute vectorsum of n-elements
- Write output array Z to FIFO
- With FIFOs we only need to store a small amount of the local data
- Does not require memory addressing
- Does not require start/done synchronization signals

```
void vectorsum ( int *X, int *Y, int *Z, int n )
{
    for ( int i = 0; i < n; i++ )
    {
        Z[i] = X[i] + Y[i];
    }
}

int main()
{
    int X[64],Y[64],Z[64];
    for ( int i=0; i < 64; i++ ) {
        X[i] = rand();
        Y[i] = rand();
    }
    vectorsum( X, Y, Z, 64 );
    :
}
```

X and Y are written  
and read in the same  
order.



# STREAMING VECTORSUM

```
module vectorsum #(
    parameter DATA_WIDTH = 32
) (
    input logic clock,
    input logic reset,
    input logic [DATA_WIDTH-1:0] x_dout,
    input logic x_empty,
    output logic x_rd_en,
    input logic [DATA_WIDTH-1:0] y_dout,
    input logic y_empty,
    output logic y_rd_en,
    output logic [DATA_WIDTH-1:0] z_din,
    input logic z_full,
    output logic z_wr_en
);

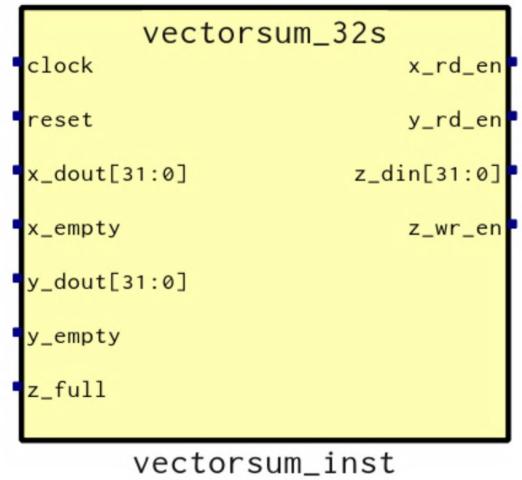
typedef enum logic [0:0] {s0, s1} state_t;
state_t state, state_c;
logic [DATA_WIDTH-1:0] sum, sum_c;

always_ff @(posedge clock or posedge reset) begin
    if (reset) begin
        state <= s0;
        sum <= '0;
    end else begin
        state <= state_c;
        sum <= sum_c;
    end
end

```

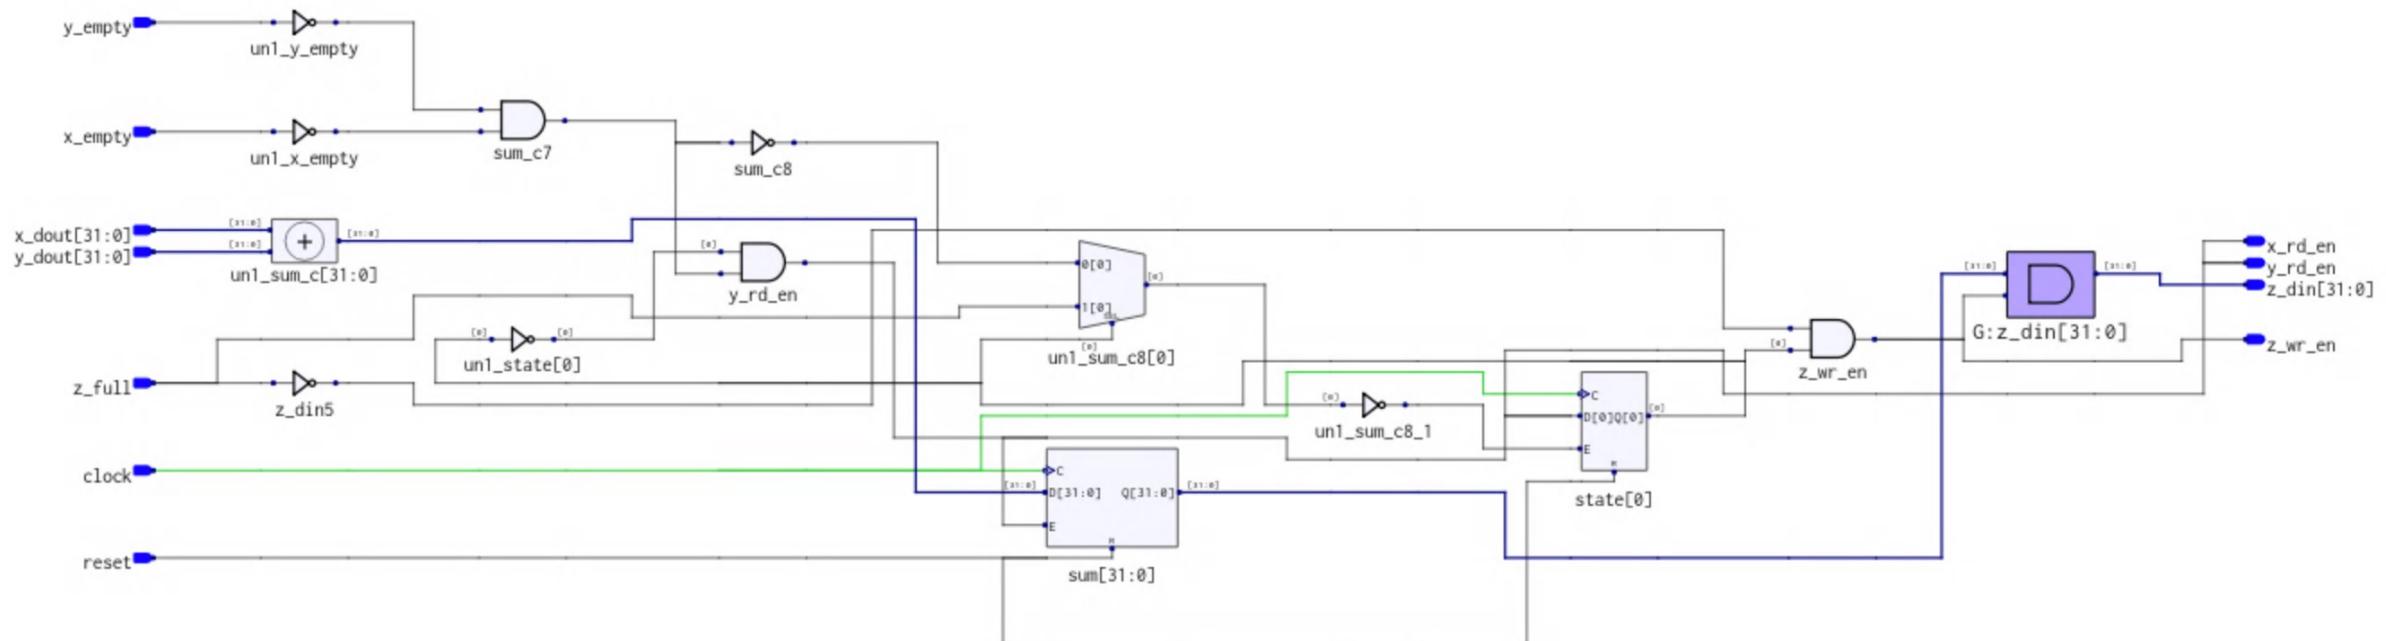
```
    always_comb begin
        z_din = 'b0;
        z_wr_en = 1'b0;
        x_rd_en = 1'b0;
        y_rd_en = 1'b0;
        sum_c = sum;
        state_c = state;
    end
    case (state)
        s0: begin
            if (x_empty == 1'b0 && y_empty == 1'b0) begin
                sum_c = $signed(y_dout) + $signed(x_dout);
                x_rd_en = 1'b1;
                y_rd_en = 1'b1;
                state_c = s1;
            end
        end
        s1: begin
            if (z_full == 1'b0) begin
                z_din = sum;
                z_wr_en = 1'b1;
                state_c = s0;
            end
        end
    endcase
end
endmodule

```



# VECTORSUM TECHNOLOGY VIEW

- Note: The vectorsum architecture is significantly simplified without the need for any memory address logic



# VECTORSUM - WRAPPER

```
module vectorsum_top #(
    parameter DATA_WIDTH = 32,
    parameter FIFO_BUFFER_SIZE = 32)
(
    input logic clock,
    input logic reset,
    output logic x_full,
    input logic x_wr_en,
    input logic [DATA_WIDTH-1:0] x_din,
    output logic y_full,
    input logic y_wr_en,
    input logic [DATA_WIDTH-1:0] y_din,
    input logic z_rd_en,
    output logic z_empty,
    output logic [DATA_WIDTH-1:0] z_dout
);

logic [DATA_WIDTH-1:0] x_dout, y_dout, z_din;
logic x_empty, y_empty, z_full;
logic x_rd_en, y_rd_en, z_wr_en;

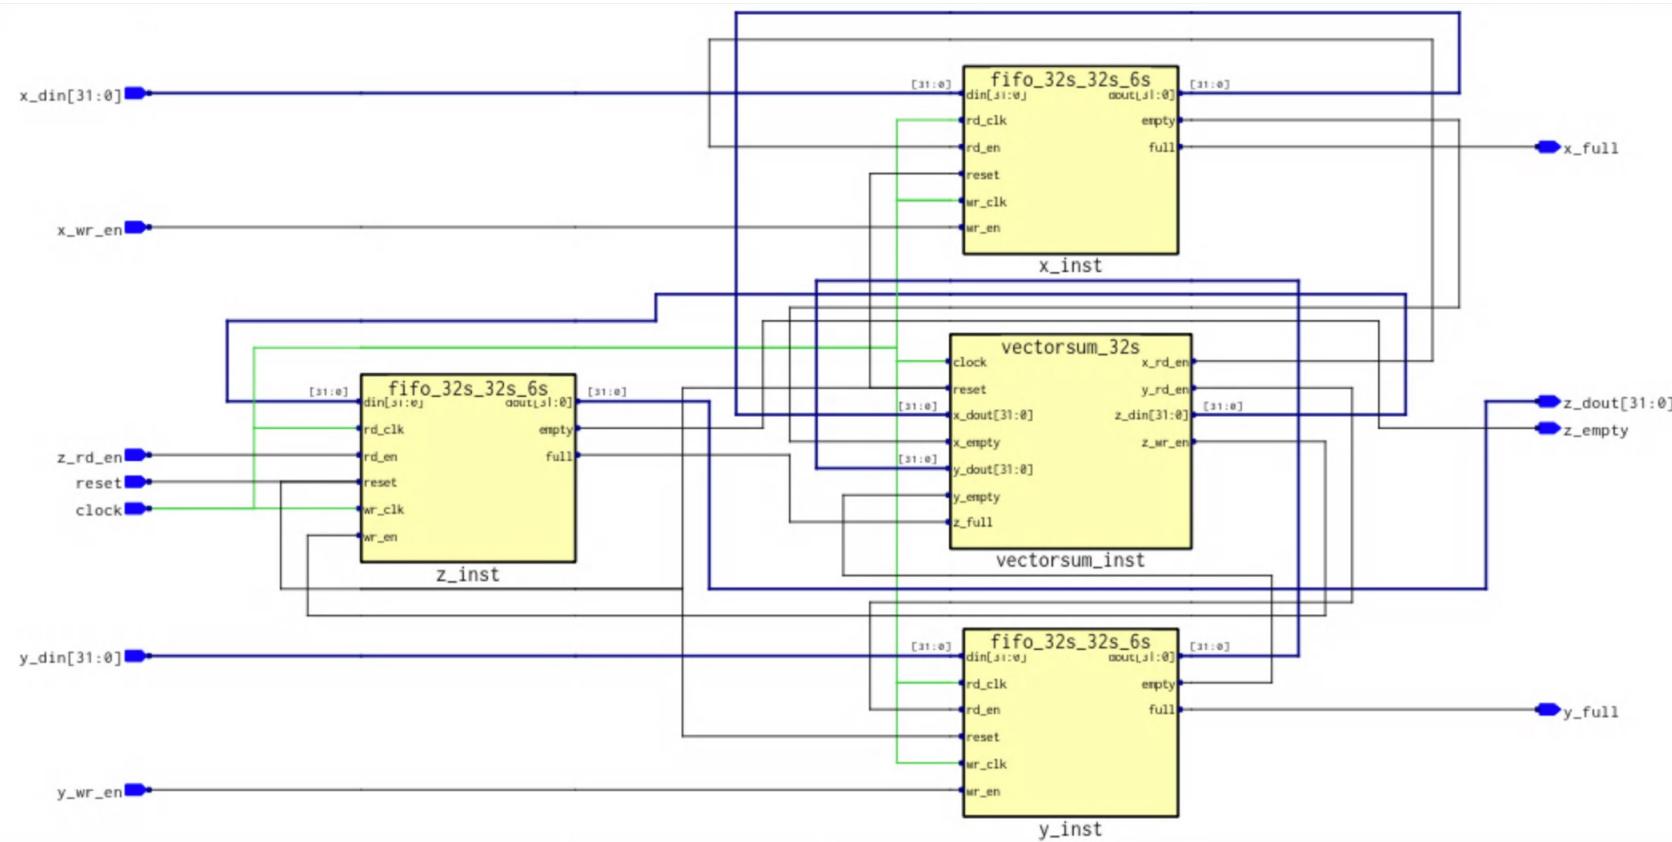
    vectorsum #(           fifo #(
        .DATA_WIDTH(DATA_WIDTH),
        .FIFO_BUFFER_SIZE(FIFO_BUFFER_SIZE),
        .FIFO_DATA_WIDTH(DATA_WIDTH)
    ) vectorsum_inst (
        .clock(clock),
        .reset(reset),
        .x_dout(x_dout),
        .x_rd_en(x_rd_en),
        .x_empty(x_empty),
        .y_dout(y_dout),
        .y_rd_en(y_rd_en),
        .y_empty(y_empty),
        .z_din(z_din),
        .z_full(z_full),
        .z_wr_en(z_wr_en)
    );

    fifo #(
        .FIFO_BUFFER_SIZE(FIFO_BUFFER_SIZE),
        .FIFO_DATA_WIDTH(DATA_WIDTH)
    ) x_inst (
        .reset(reset),
        .wr_clk(clock),
        .wr_en(x_wr_en),
        .din(x_din),
        .full(x_full),
        .rd_clk(clock),
        .rd_en(x_rd_en),
        .dout(x_dout),
        .empty(x_empty)
    );

    fifo #(
        .FIFO_BUFFER_SIZE(FIFO_BUFFER_SIZE),
        .FIFO_DATA_WIDTH(DATA_WIDTH)
    ) y_inst (
        .reset(reset),
        .wr_clk(clock),
        .wr_en(y_wr_en),
        .din(y_din),
        .full(y_full),
        .rd_clk(clock),
        .rd_en(y_rd_en),
        .dout(y_dout),
        .empty(y_empty)
    );

    fifo #(
        .FIFO_BUFFER_SIZE(FIFO_BUFFER_SIZE),
        .FIFO_DATA_WIDTH(DATA_WIDTH)
    ) z_inst (
        .reset(reset),
        .wr_clk(clock),
        .wr_en(z_wr_en),
        .din(z_din),
        .full(z_full),
        .rd_clk(clock),
        .rd_en(z_rd_en),
        .dout(z_dout),
        .empty(z_empty)
    );
endmodule
```

# VECTORSUM TOP-LEVEL TECHNOLOGY VIEW



# TESTBENCH

```
'timescale 1 ns / 1 ns

module vectorsum_tb ();
localparam string X_NAME = "x.txt";
localparam string Y_NAME = "y.txt";
localparam string Z_NAME = "z.txt";
localparam DATA_WIDTH = 32;
localparam ADDR_WIDTH = 10;
localparam VECTOR_SIZE = 64;
localparam FIFO_BUFFER_SIZE = 8;
localparam CLOCK_PERIOD = 10;

logic clock = 1'b0;
logic reset = 1'b0;

logic [DATA_WIDTH-1:0] x_din;
logic x_wr_en, x_full;
logic [DATA_WIDTH-1:0] y_din;
logic y_wr_en, y_full;
logic [DATA_WIDTH-1:0] z_dout;
logic z_rd_en, z_empty;

logic x_write_done = '0;
logic y_write_done = '0;
logic z_read_done = '0;
integer z_errors = '0;

vectorsum_top #(           // tb process
    .DATA_WIDTH(DATA_WIDTH),
    .FIFO_BUFFER_SIZE(FIFO_BUFFER_SIZE)
) vectorsum_top_inst (
    .clock(clock),
    .reset(reset),
    .x_full(x_full),
    .x_wr_en(x_wr_en),
    .x_din(x_din),
    .y_full(y_full),
    .y_wr_en(y_wr_en),
    .y_din(y_din),
    .z_rd_en(z_rd_en),
    .z_empty(z_empty),
    .z_dout(z_dout)
);

// clock process
always begin
    #(CLOCK_PERIOD/2) clock = 1'b1;
    #(CLOCK_PERIOD/2) clock = 1'b0;
end

// reset process
initial begin
    #(CLOCK_PERIOD) reset = 1'b1;
    #(CLOCK_PERIOD) reset = 1'b0;
end

initial begin
    time start_time, end_time;
    @(negedge reset);
    @(posedge clock);
    start_time = $time;
    $display("@ %0t: Beginning simulation...", start_time);
    wait(z_read_done);

    end_time = $time;
    $display("@ %0t: Simulation completed.", end_time);
    $display("Total simulation cycle count: %0d",
        (end_time-start_time)/CLOCK_PERIOD);
    $display("Total error count: %0d", z_errors);

    $stop;
end
```

# TESTBENCH: READ FILE INTO FIFO

```
initial begin : x_write
    integer fd, x=0, count;
    x_din = '0;
    x_wr_en = 1'b0;

    @(negedge reset);
    $display("@ %0t: Loading file %s...", $time, X_NAME);
    fd = $fopen(X_NAME, "r");

    while ( x < VECTOR_SIZE ) begin
        @(negedge clock);
        x_wr_en = 1'b0;
        if (x_full == 1'b0) begin
            count = $fscanf(fd, "%h", x_din);
            x_wr_en = 1'b1;
            x++;
        end
    end

    @(posedge clock);
    x_wr_en = 1'b0;
    $fclose(fd);
    x_write_done = 1'b1;
end

initial begin : y_write
    integer fd, y=0, count;
    y_din = '0;
    y_wr_en = 1'b0;
    @(negedge reset);
    $display("@ %0t: Loading file %s...", $time, Y_NAME);
    fd = $fopen(Y_NAME, "r");

    while ( y < VECTOR_SIZE ) begin
        @(negedge clock);
        y_wr_en = 1'b0;
        if (y_full == 1'b0) begin
            count = $fscanf(fd, "%h", y_din);
            y_wr_en = 1'b1;
            y++;
        end
    end

    @(posedge clock);
    y_wr_en = 1'b0;
    $fclose(fd);
    y_write_done = 1'b1;
end
```

# TESTBENCH: READ FIFO AND COMPARE OUTPUT

```
initial begin : z_write
    integer fd, z=0, count;
    logic [DATA_WIDTH-1:0] z_data_cmp;
    z_rd_en = 1'b0;
    z_data_cmp = '0;

    @(negedge reset);
    @(negedge clock);

    $display("@ %0t: Comparing file %s...", $time, Z_NAME);
    fd = $fopen(Z_NAME, "r");

    while ( z < VECTOR_SIZE ) begin
        @(negedge clock);
        z_rd_en = 1'b0;
        if (z_empty == 1'b0) begin
            z_rd_en = 1'b1;
            count = $fscanf(fd, "%h", z_data_cmp);
            if (z_dout != z_data_cmp) begin
                z_errors++;
                $display("@ %0t: %s(%0d): ERROR: %h != %h at address 0x%h.", $time, Z_NAME, z+1, z_dout, z_data_cmp, z);
            end
            z++;
        end
        z++;
    end
    z_rd_en = 1'b0;
    z_read_done = 1'b1;
    $fclose(fd);
end
```

# SIMULATION SCRIPT

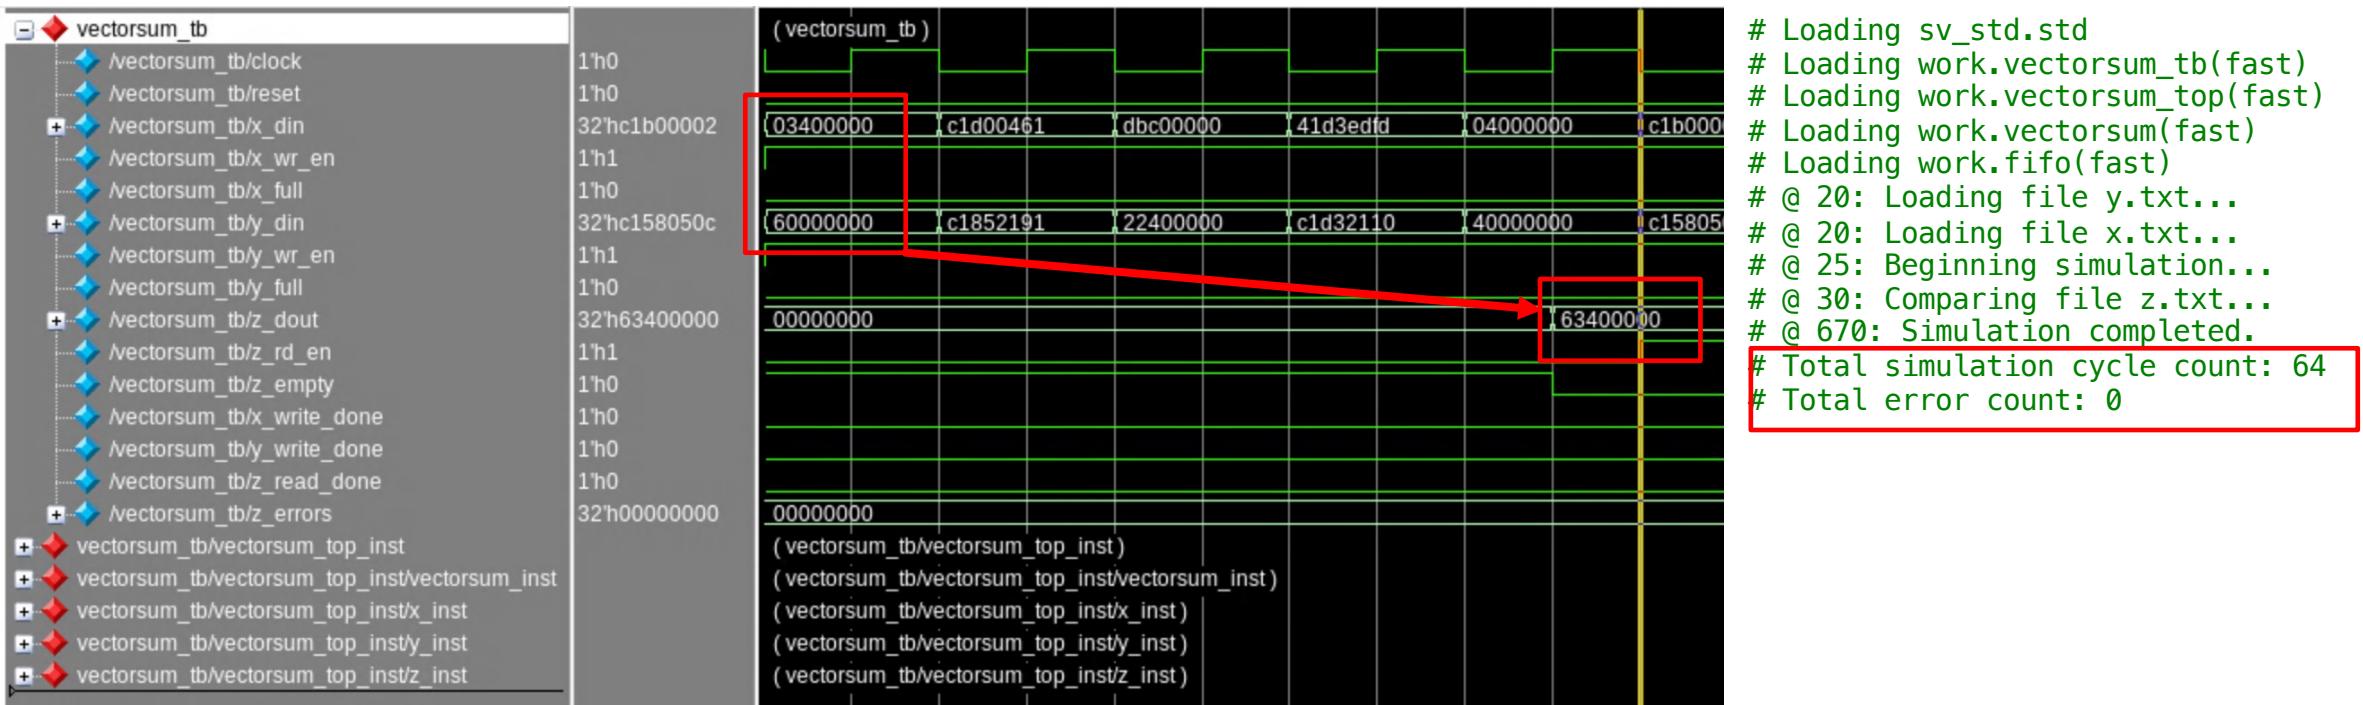
```
setenv LMC_TIMEUNIT -9
vlib work
vmap work work
vlog -work work "../sv/fifo.sv"
vlog -work work "../sv/vectorsum.sv"
vlog -work work "../sv/vectorsum_top.sv"
vlog -work work "../sv/vectorsum_tb.sv"
vsim -classdebug -voptargs=+acc +notimingchecks -L work work.vectorsum_tb -wlf vectorsum.wlf

add wave -noupdate -group vectorsum_tb
add wave -noupdate -group vectorsum_tb -radix hexadecimal /vectorsum_tb/*
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst -radix hexadecimal /vectorsum_tb/vectorsum_top_inst/*
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/vectorsum_inst
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/vectorsum_inst -radix hexadecimal /vectorsum_tb/vectorsum_top_inst/vectorsum_inst/*
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/x_inst
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/x_inst -radix hexadecimal /vectorsum_tb/vectorsum_top_inst/x_inst/*
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/y_inst
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/y_inst -radix hexadecimal /vectorsum_tb/vectorsum_top_inst/y_inst/*
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/z_inst
add wave -noupdate -group vectorsum_tb/vectorsum_top_inst/z_inst -radix hexadecimal /vectorsum_tb/vectorsum_top_inst/z_inst/*

run -all
```

# MODELSIM SIMULATION

- Validation of errors and cycle count



NEXT...

- HW #3: Streaming architecture with Motion Detection algorithm